the globus alliance
www.globus.org

OGSA-DAI Today

GridWorld 2006, Washington DC

11 September 2006

epcc

Outline

- What is it?
- "Let us out"
  - Exposing data to clients – the server's perspective
- "Let us in"
  - Getting to the data – the client's perspective
- "More, more more…"
  - Extending OGSA-DAI

OGSA-DAI in a nutshell

- An *extensible framework* for data access and integration
- Expose heterogeneous data resources to a grid through web services
- Interact with data resources
  - Queries and updates
  - Data transformation / compression
  - Data delivery
  - Application-specific functionality
- A base for higher-level services
  - Federation, mining, visualisation,…

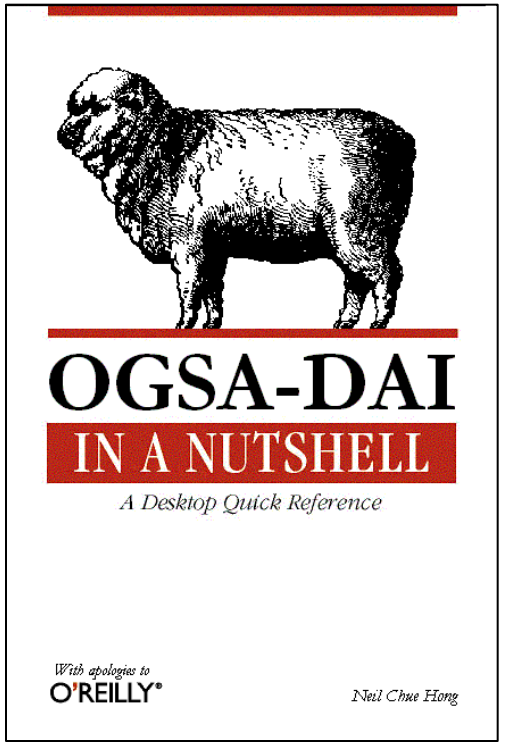

OGSA-DAI motivation

- Entering an age of data
  - Data Explosion
    - CERN LHC will generate 1GB/s = 10PB/y
    - VLBA (NRAO) generates 1GB/s today
    - Pixar generate 100 TB/movie
  - Storage getting cheaper
- Data stored in many different ways
  - Relational databases
  - XML databases
  - Text and binary files
- Need ways to facilitate
  - Data discovery
  - Data access
  - Data integration
- Empower e-Business and e-Science
  - The grid is a vehicle for achieving this

Data resources

- Relational
  - MySQL
  - Microsoft SQL Server
  - Oracle
  - IBM DB2
  - PostGres
  - HSQL
- XML
  - eXist
  - Xindice
- File system
  - SwissPROT
  - OMIM
  - Text
  - Binary

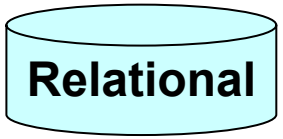

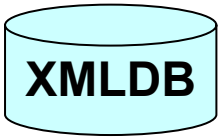

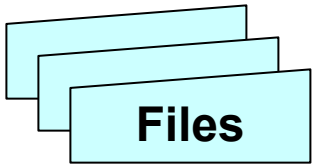

Data resource accessors

- Interfaces between data resources and OGSA-DAI
- Relational
  - JDBC drivers
  - `java.sql.*`
- XML
  - XMLDB API and compliant drivers
  - `org.xmldb.api.*`
- File system
  - Java file and directory utilities
  - `java.io.*`

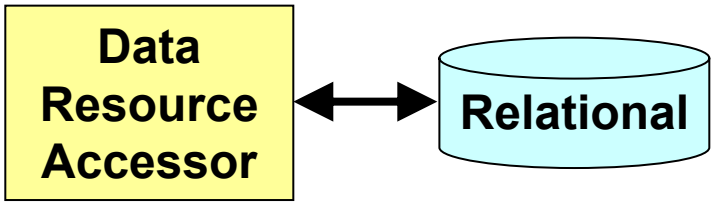

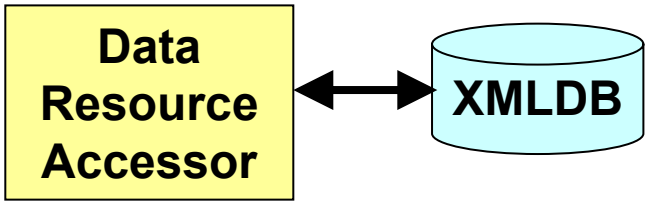

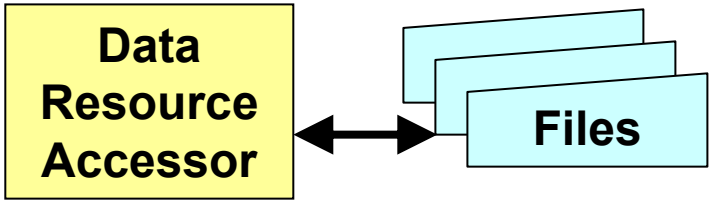

Data service resources

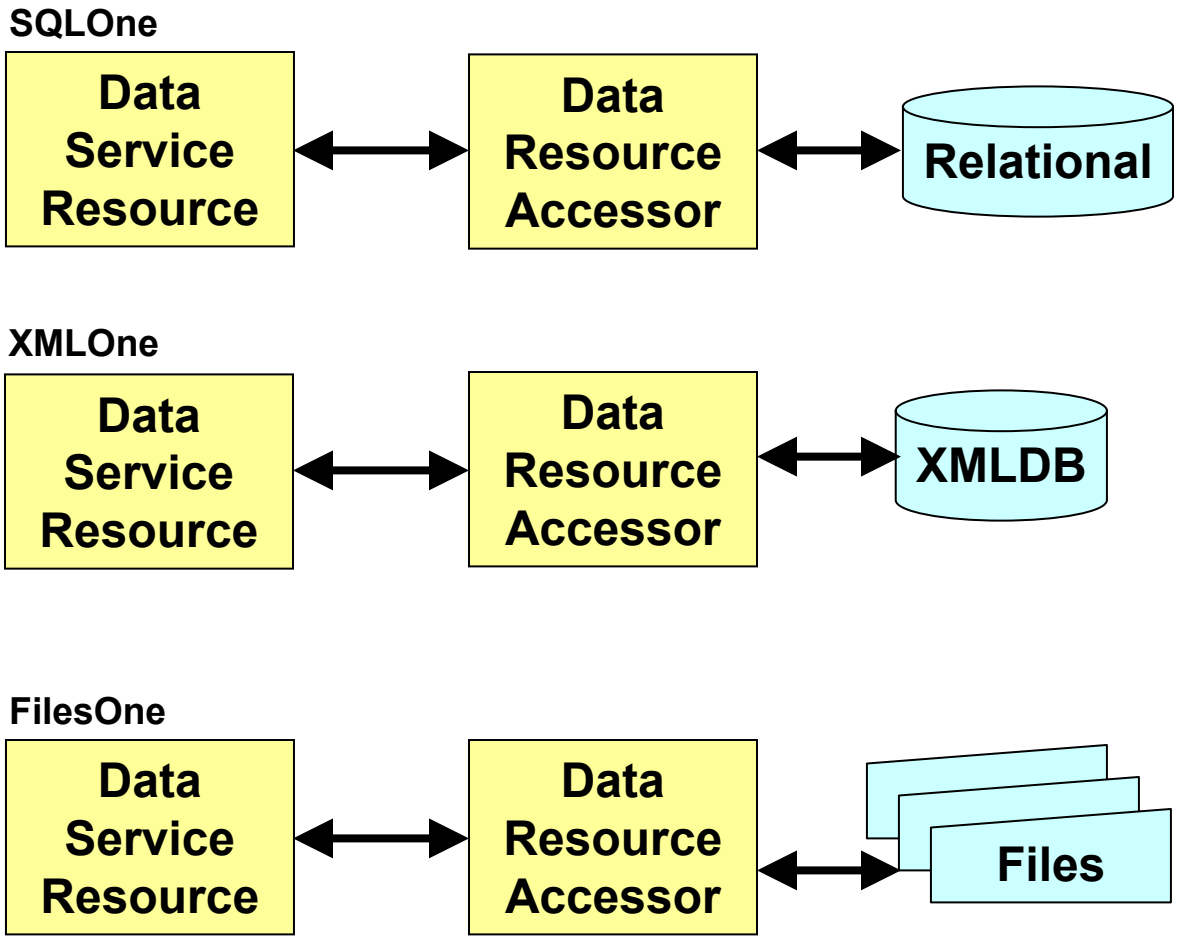

Data service resources

- OGSA-DAI's core functionality
- Manages
  - Access to a data resource via a data resource accessor
  - Execution of data-related activities
  - Data caching and streaming of data to and from clients
  - Creation, access and termination of sessions
- Exposes data service resource properties
  - Information about a data resource
  - Information about supported activities
  - Information about current requests

Requests and responses

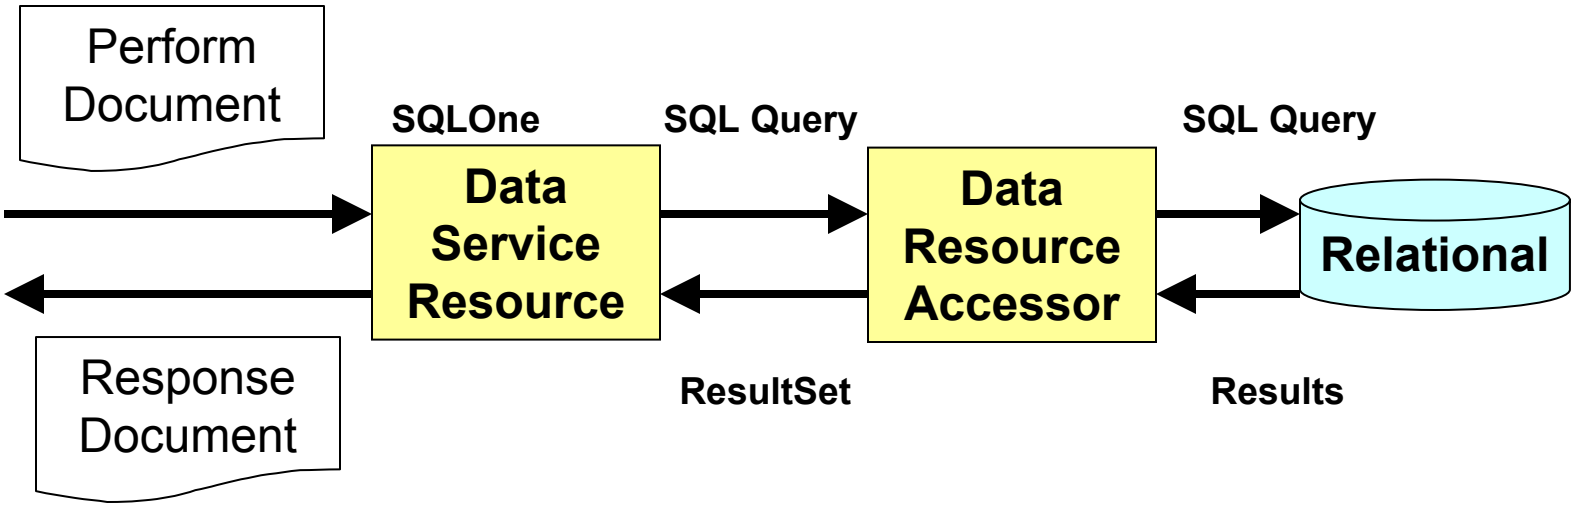

Requests and responses

- Request
  - A connected collection of activities that the data resource executes
  - Flow control – sequential or parallel execution of activities
  - XML perform document submitted by a client
- Activity
  - An individual data-related operation
  - 0 or more inputs and 0 or more outputs
- Response
  - Status of execution of a request possibly with result data
  - XML response document returned to a client
- OGSA-DAI engine
  - Parses requests, executes activities, builds responses

Activities

- Relational
  - SQL query, update, stored procedure, bulk load, extract logical and physical schema
  - Convert ResultSet to WebRowSet, ResultSet to CSV, ResultSet to bytes, relational database schema to XML
  - Project ResultSet or WebRowSet onto a column
  - Extract bytes from ResultSet
- XMLDB
  - Resource and collection management, XPath, XQuery, XUpdate, bulk load
- Files
  - List directory, create, read, write and update files
  - Index files, search indexed files

Activities

- Transformation and Compression
  - GZIP compression, ZIP archive
  - XSLT
  - Project CSV data onto a column
  - Distribute numerical data onto spaces
  - Create random sample of data
- Delivery
  - From and to URLs, files, GridFTP, remote data service resources, SOAP attachments
  - To servlets, SMTP, resource properties
- Factory
  - Create/destroy persistent/transient data service resources
- Relational multi-resources
  - Bag and resilient queries

Data services

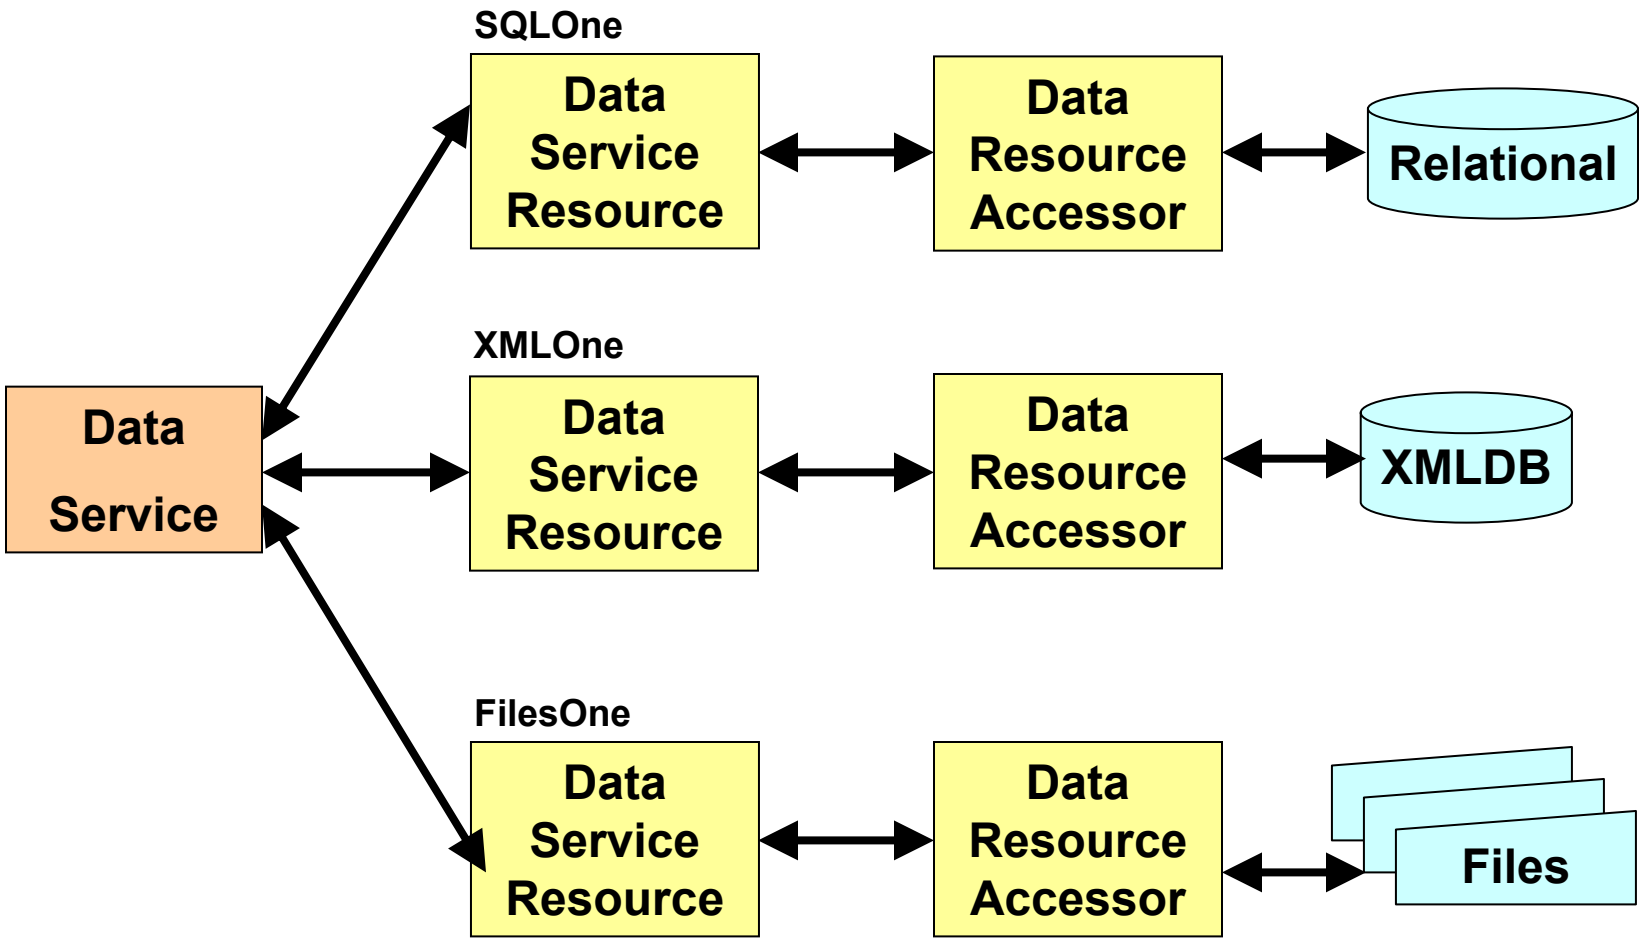

Data services

- Web services
- Expose 0..N data service resources to the outside world
- Two flavours
  - OGSA-DAI WSRF services
    - Compliant with the Web Services Resource Framework
    - Implemented using Globus Toolkit (4.0+)
  - OGSA-DAI WSI services
    - Compliant with vanilla WSDL
    - Implemented using Apache Axis (1.2.1 or 1.2RC3)

Clients

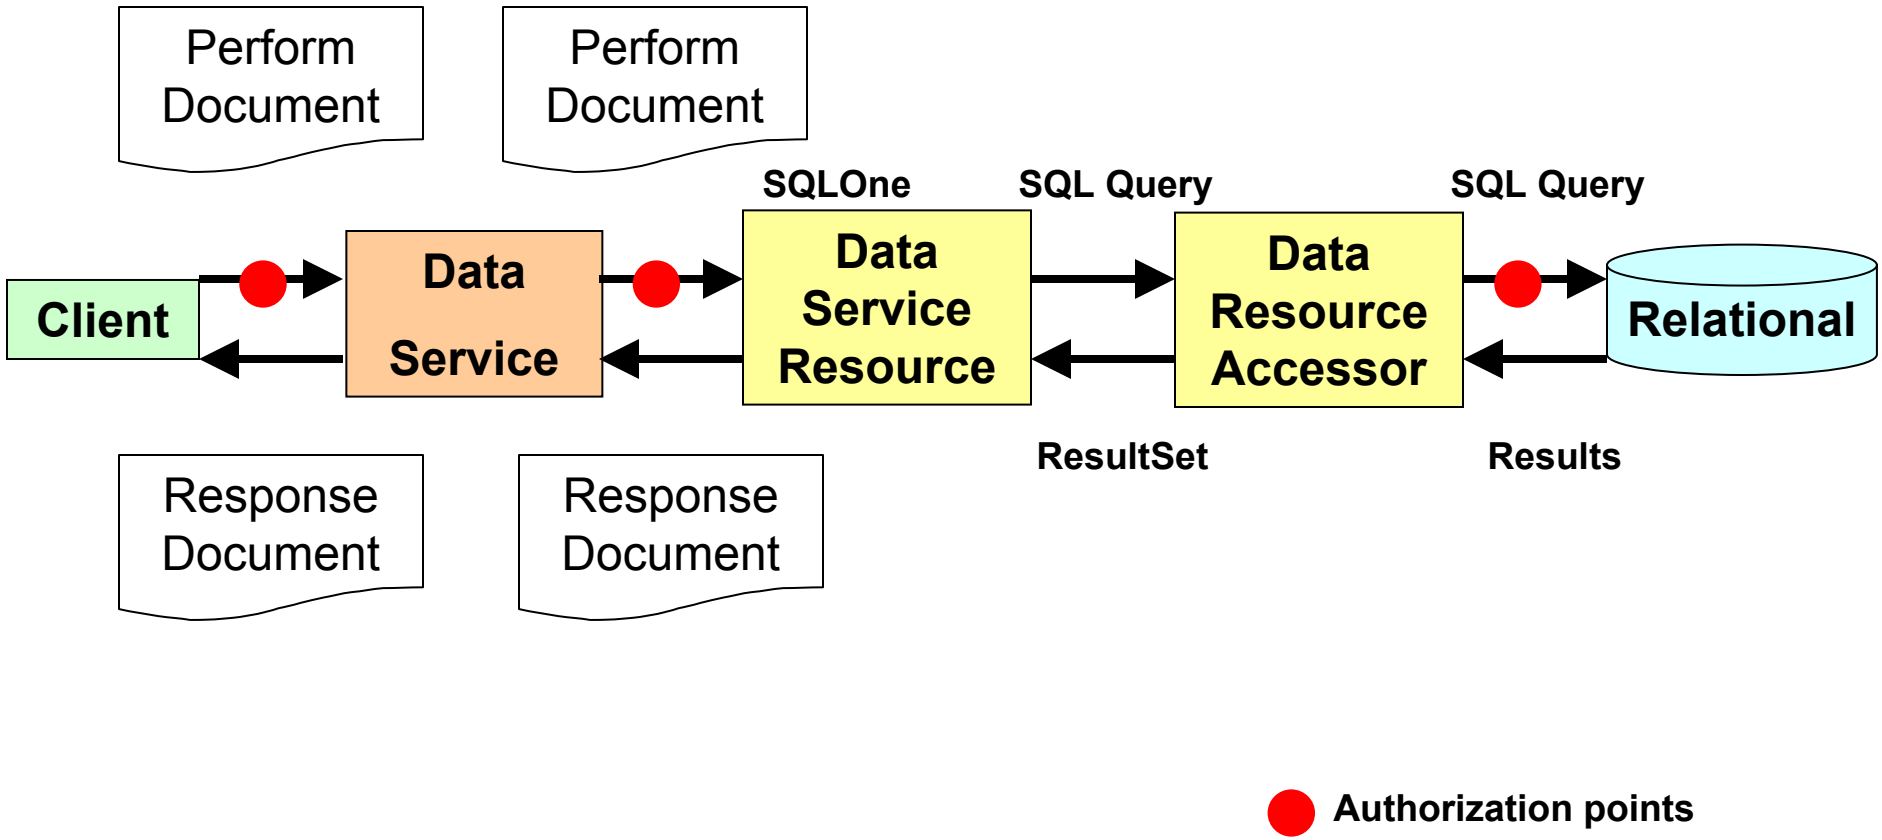
Perform Document
Perform Document
SQLOne
SQL Query
SQL Query
Client
Data Service
Data Service Resource
Data Resource Accessor
Relational
ResultSet
Results
Response Document
Response Document
Authorization points

Identifying a data service resource

http://host:port/services/wsrf/DataService

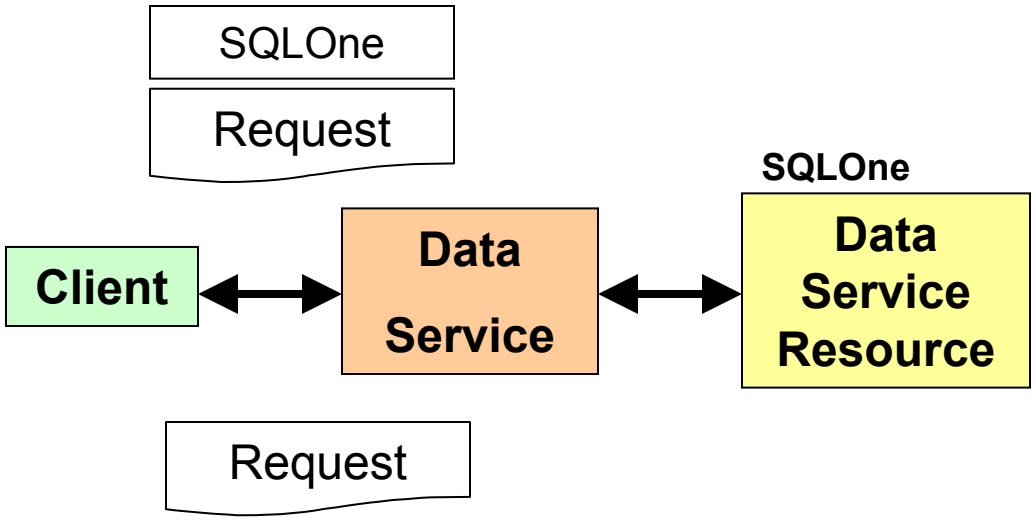

http://host:port/services/axis/DataService/DAISQLOne

Clients and the client toolkit

- Clients interact with data services via SOAP over HTTP
  - Deduce service interface from service WSDL description
  - Construct SOAP request to invoke operation
  - Parse SOAP response from service
  - Resource identification scheme must be assumed from WSDL namespace
- OGSA-DAI client toolkit:
  - Construct and submit requests in Java not XML
    - Toolkit handles SOAP request construction and response parsing
  - Renders OGSA-DAI service types transparent
  - Java abstractions of
    - Data services
    - Data service resource IDs and session IDs
    - Requests and responses
    - Activities

Relational multi-resources

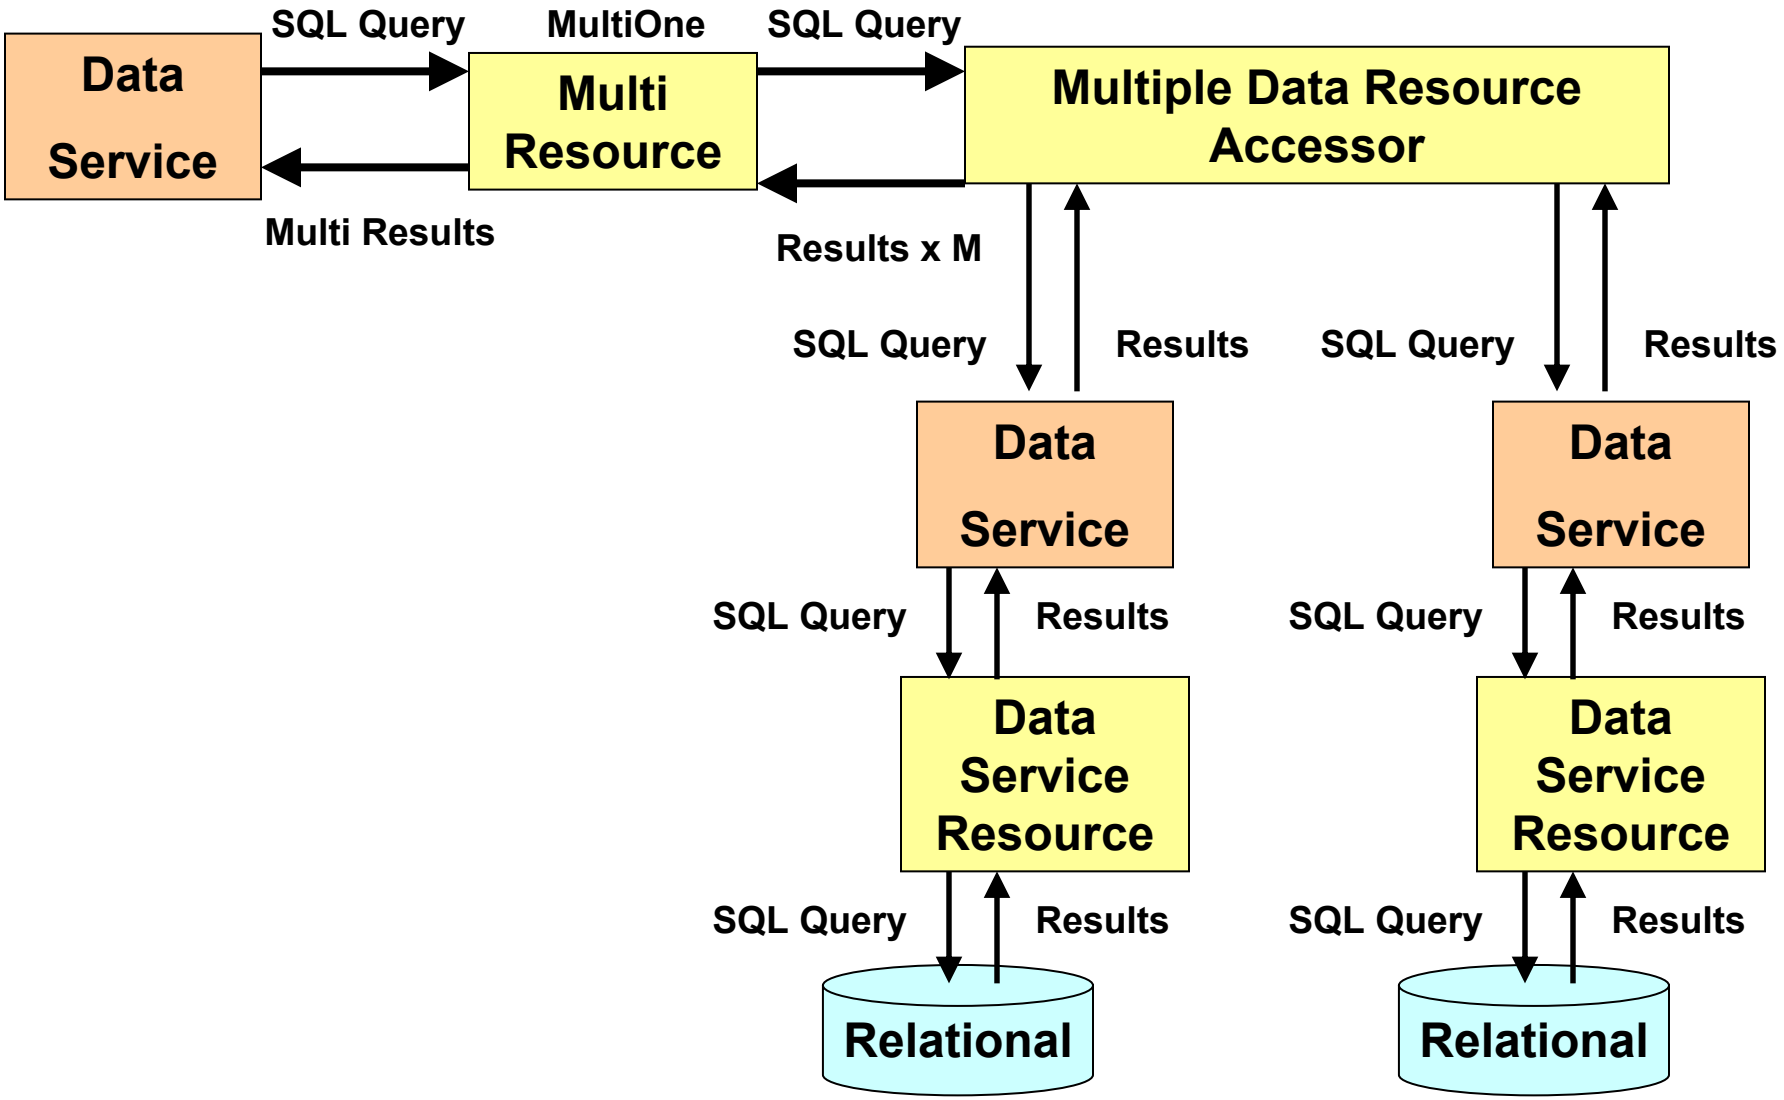

Factory resources

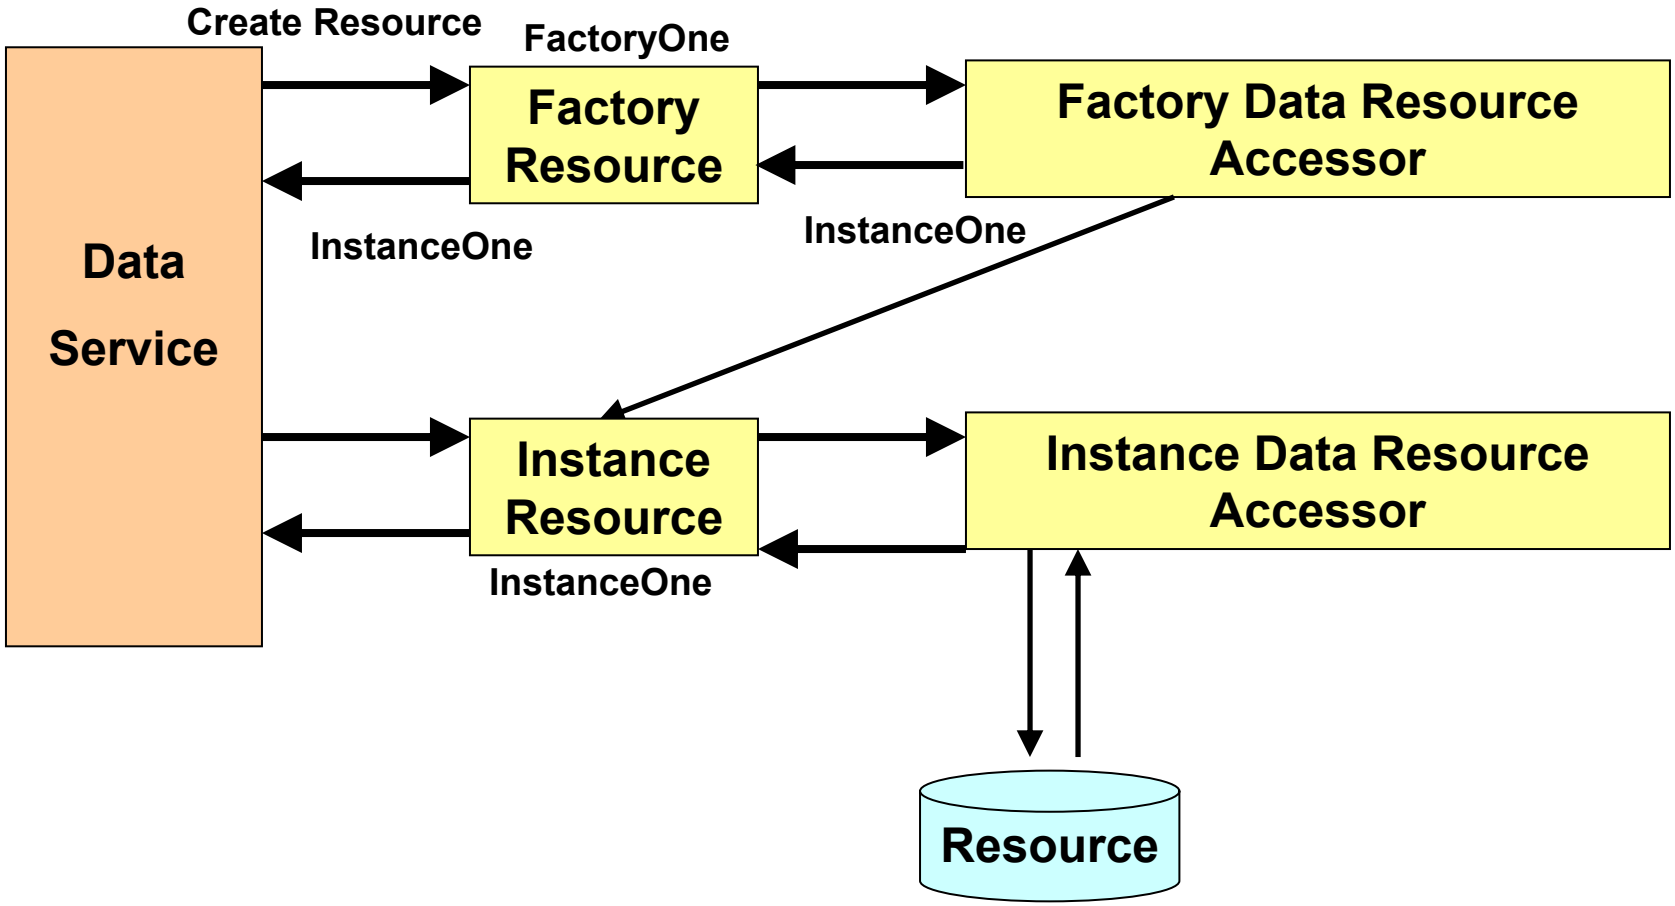

Extending OGSA-DAI

- Application-specific data resource accessors
  - Expose local or remote data resources
  - Expose virtual resources created by aggregation or integration
  - Create/destroy of persistent/transient data service resources
- Application-specific activities
  - Can be resource specific e.g query or update
  - Or generic e.g. transformation, compression, delivery, resource management, monitoring
- Application-specific authorization
  - Resource access
  - Activity execution

OGSA-DAI today - what's in a name?

- April 28th 2006
- It's not release 8
  - The lingering spectre of OGSI
- It's release 2.2
  - OGSA-DAI WSRF 2.2
    - Runs under Globus Toolkit 4.0.1 and 4.0.2
    - Patched version bundled with Globus Toolkit 4.1.0
  - OGSA-DAI WSI 2.2
    - Runs under Apache Axis 1.2.1 on Tomcat
    - Runs under Axis 1.2RC3 on Tomcat
    - Runs under OMII 3.0.1
    - Patched version bundled with OMII 3.
- http://www.ogsadai.org.uk

Resource management

- Transient data service resources
  - Exist only in memory
  - No associated configuration files on the server
  - Associated activity to create new transient resources
- Resource withdrawal
  - An activity to withdraw a data service resource
    - Optionally remove its configuration files from the server

Relational multi-resources

- A data service resource that aggregates multiple relational data service resources
- Aggregated resources can be local or remote
- Associated activities
  - Submit a query to every aggregated resource and return a bag of the results
  - Submit a query to every aggregated resource and return the results from the first one that completes
  - Remove duplicate rows from a bag of the results

Activities...

- Data conversion activities
  - Convert ResultSet to WebRowSet
  - Convert ResultSet to CSV
  - Convert BLOB from ResultSet column into bytes
- Relational meta-data activities
  - Retrieve logical database schema
  - Convert logical database schema to XML
  - Retrieve physical database schema
- XMLDB activities
  - XQuery

...and some more

- Projection and transformation activities
  - Remove duplicate rows from a WebRowSet
  - Project a ResultSet onto a column name or index
  - Project a WebRowSet onto a column name or index
  - Project CSV values onto a column index
  - Distribute numeric values onto spaces
  - Generate a random sample of input data
  - Write a stream of bytes to a temporary file and output a reference to this file
- Delivery activities
  - Throw data away
  - Write data to a resource property
  - Deliver data to a SOAP attachment

Then there's…

- eXist data resources and XQuery
- BLOBs
  - Improved support for BLOBs in SQL query and update activities
  - Activity to dump BLOBs into temporary files server-side
- Security
  - Resource and activity authorization
  - GSI Secure Conversation message-level security for inter-service communications using data transport
    - OGSA-DAI WSRF only
- Revamped logging, exceptions and internationalization
- Usage
  - Publication of an initial set of usage scenarios and best practice

…and…

- Bundled third-party JARs
  - OGSA-DAI WSI
    - JARs required to compile OGSA-DAI source distribution JARs required to run OGSA-DAI clients
  - OGSA-DAI WSRF
    - Non-Globus Toolkit JARs required to compile OGSA-DAI source distribution
    - JARs required to run OGSA-DAI clients
- Benchmarking and performance
  - WebRowSet – up to 35% better than in 2.0
  - resultSetToCSV activity – CSV yields up to 65% improvement compared to WebRowSet
  - Binary data transfer
    - Takes 25% of the time to transfer a binary 8MB file using SOAP attachments instead of in SOAP body
    - Improvements due to smaller data size and smaller SOAP messages
    - Limited by I/O performance rather than CPU
  - Numerous other improvements

Summary

- OGSA-DAI is an extensible framework for building data access and integration applications
- The OGSA-DAI layer cake
  - Data
  - Data resource accessors
  - Data service resources
  - Data services
  - Clients
- Extending OGSA-DAI
  - Data resource accessors
  - Activities
  - Authorization

Further information

- The OGSA-DAI project site
  - http://www.ogsadai.org.uk
- The DAIS-WG site
  - http://forge.gridforum.org/projects/dais-wg
- OGSA-DAI users mailing list
  - users@ogsadai.org.uk
  - General discussion on OGSA-DAI, data and the grid
- Formal support for OGSA-DAI releases
  - http://www.ogsadai.org.uk/support
  - support@ogsadai.org.uk
- OGSA-DAI training courses